

# RFM Analysis in Python

RFM analysis is scoring our customers based on their Recency, Frequency and Monetary values. - Recency: How recently a customer made a purchase. - Frequency: How often customers make a purchase. - Monetary Value: How much money a customer spends on purchases.

```
In [38]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from datetime import datetime
import datetime as dt
import warnings
warnings.filterwarnings('ignore')
```

## Loading data

```
In [39]: data=pd.read_csv('sales_data.csv')
data
```

```
Out[39]:
```

	OrderNumber	Sales Channel	WarehouseCode	ProcuredDate	OrderDate	ShipDate	DeliveryDate	CurrencyCode	_SalesTeamID
0	SO - 000101	In-Store	WARE-UHY1004	12/31/2017	5/31/2018	6/14/2018	6/19/2018	USD	6
1	SO - 000102	Online	WARE-NMK1003	12/31/2017	5/31/2018	6/22/2018	7/2/2018	USD	14
2	SO - 000103	Distributor	WARE-UHY1004	12/31/2017	5/31/2018	6/21/2018	7/1/2018	USD	21
3	SO - 000104	Wholesale	WARE-NMK1003	12/31/2017	5/31/2018	6/2/2018	6/7/2018	USD	28
4	SO - 000105	Distributor	WARE-NMK1003	4/10/2018	5/31/2018	6/16/2018	6/26/2018	USD	22
...	...	...	...	...	...	...	...	...	...
7986	SO - 0008087	In-Store	WARE-MKL1006	9/26/2020	12/30/2020	1/7/2021	1/14/2021	USD	9
7987	SO - 0008088	Online	WARE-NMK1003	9/26/2020	12/30/2020	1/2/2021	1/4/2021	USD	14
7988	SO - 0008089	Online	WARE-UHY1004	9/26/2020	12/30/2020	1/23/2021	1/26/2021	USD	14
7989	SO - 0008090	Online	WARE-NMK1003	9/26/2020	12/30/2020	1/20/2021	1/25/2021	USD	20
7990	SO - 0008091	In-Store	WARE-UHY1004	9/26/2020	12/30/2020	1/13/2021	1/19/2021	USD	6

7991 rows × 16 columns

## Understanding Data

```
In [36]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7991 entries, 0 to 7990
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   OrderNumber           7991 non-null   object
1   Sales Channel         7991 non-null   object
2   WarehouseCode         7991 non-null   object
3   ProcuredDate          7991 non-null   object
4   OrderDate             7991 non-null   object
5   ShipDate              7991 non-null   object
6   DeliveryDate          7991 non-null   object
7   CurrencyCode          7991 non-null   object
8   _SalesTeamID          7991 non-null   int64
9   _CustomerID           7991 non-null   int64
10  _StoreID              7991 non-null   int64
11  _ProductID            7991 non-null   int64
12  Order Quantity        7991 non-null   int64
13  Discount Applied      7991 non-null   float64
14  Unit Price            7991 non-null   float64
15  Unit Cost             7991 non-null   float64
dtypes: float64(3), int64(5), object(8)
memory usage: 999.0+ KB
```

```
In [40]: data.shape
```

```
Out[40]: (7991, 16)
```

```
In [41]: data.columns
```

```
Out[41]: Index(['OrderNumber', 'Sales Channel', 'WarehouseCode', 'ProcuredDate',
        'OrderDate', 'ShipDate', 'DeliveryDate', 'CurrencyCode', '_SalesTeamID',
        '_CustomerID', '_StoreID', '_ProductID', 'Order Quantity',
        'Discount Applied', 'Unit Price', 'Unit Cost'],
        dtype='object')
```

```
In [42]: data.describe()
```

```
Out[42]:
```

	_SalesTeamID	_CustomerID	_StoreID	_ProductID	Order Quantity	Discount Applied	Unit Price	Unit Cost
count	7991.000000	7991.000000	7991.000000	7991.000000	7991.000000	7991.000000	7991.000000	7991.000000
mean	14.384307	25.457014	183.850081	23.771743	4.525341	0.114394	2284.536504	1431.911054
std	7.986086	14.414883	105.903946	13.526545	2.312631	0.085570	1673.096364	1112.413043
min	1.000000	1.000000	1.000000	1.000000	1.000000	0.050000	167.500000	68.675000
25%	8.000000	13.000000	91.000000	12.000000	3.000000	0.050000	1031.800000	606.115500
50%	14.000000	25.000000	183.000000	24.000000	5.000000	0.075000	1849.200000	1080.576000
75%	21.000000	38.000000	276.000000	36.000000	7.000000	0.150000	3611.300000	2040.250500
max	28.000000	50.000000	367.000000	47.000000	8.000000	0.400000	6566.000000	5498.556000

## Data preparation

```
In [43]: #checking null values:
data.isnull().sum()
```

```
Out[43]: OrderNumber      0
Sales Channel      0
WarehouseCode      0
ProcuredDate      0
OrderDate      0
ShipDate      0
DeliveryDate      0
CurrencyCode      0
_SalesTeamID      0
_CustomerID      0
_StoreID      0
_ProductID      0
Order Quantity      0
Discount Applied      0
Unit Price      0
Unit Cost      0
dtype: int64
```

```
In [44]: #Dropping Duplicated Values
data = data.drop_duplicates()
```

```
In [45]: #how many of each product?
data["WarehouseCode"].value_counts().head()
```

```
Out[45]: WARE-NMK1003    2505
WARE-PUJ1005    1451
WARE-UHY1004    1265
WARE-XYS1001    1222
WARE-MKL1006     857
Name: WarehouseCode, dtype: int64
```

```
In [46]: #What are the most expensive products?
data.sort_values("Unit Price", ascending = False).head()
```

```
Out[46]:
```

	OrderNumber	Sales Channel	WarehouseCode	ProcuredDate	OrderDate	ShipDate	DeliveryDate	CurrencyCode	_SalesTeamID
1776	SO - 0001877	Distributor	WARE-PUJ1005	10/27/2018	12/26/2018	1/3/2019	1/13/2019	USD	23
3963	SO - 0004064	In-Store	WARE-XYS1001	5/15/2019	9/12/2019	9/23/2019	9/29/2019	USD	11
4350	SO - 0004451	Wholesale	WARE-NMK1003	8/23/2019	10/31/2019	11/10/2019	11/17/2019	USD	28
3603	SO - 0003704	Online	WARE-UHY1004	2/4/2019	7/29/2019	8/18/2019	8/27/2019	USD	15
5265	SO - 0005366	Distributor	WARE-NMK1003	12/1/2019	2/16/2020	2/26/2020	2/28/2020	USD	25

```
In [47]: #Creating 'Total Price' Column
data['Total Price'] = data['Unit Price']*data['Order Quantity']
data['Total Price']
```

```
Out[47]: 0      9815.5
        1     11818.8
        2     1775.5
        3    18599.2
        4    14579.2
        ...
        7986     234.5
        7987    19215.6
        7988    19128.5
        7989     8576.0
        7990    11055.0
Name: Total_Price, Length: 7991, dtype: float64
```

## Creating RFM Dataframe

```
In [48]: #converting 'OrderDate' column to datetime format
data['OrderDate'] = pd.to_datetime(data['OrderDate'])
```

```
In [49]: #last purchase date:
data['OrderDate'].max()
```

```
Out[49]: Timestamp('2020-12-30 00:00:00')
```

```
In [50]: #date for analysis :
today = dt.datetime(2020,12,30)
```

```
In [51]: today
```

```
Out[51]: datetime.datetime(2020, 12, 30, 0, 0)
```

```
In [54]: rfm_data = data.groupby('_CustomerID').agg({
    'OrderDate': lambda x: (today - x.max()).days, #For Recency, Calculate the number of days between present d
    '_CustomerID': 'count', #For Frequency, Calculate the number of orders for each customer.
    'Total_Price': 'sum' #For Monetary, Calculate sum of purchase price for each customer.
})

# Rename columns for clarity
rfm_data.rename(columns={
    'OrderDate': 'Recency',
    '_CustomerID': 'Frequency',
    'Total_Price': 'Monetary'
}, inplace=True)
```

```
In [56]: rfm_data.head()
```

```
Out[56]:
```

	Recency	Frequency	Monetary
_CustomerID			
1	7	152	1322278.5
2	7	135	1346264.5
3	8	181	1831947.5
4	3	167	1770582.2
5	28	159	1609232.8

## RFM Quartiles:

Dividing the RFM values into quartiles to categorize customers.

```
In [59]: # Calculate quartiles for Recency, Frequency, and Monetary
quartiles = rfm_data.quantile(q=[0.25, 0.5, 0.75])

# Create functions to assign R, F, and M scores based on quartiles
def assign_fm_score(x, quartiles):
    if x <= quartiles[0.25]:
        return 1
    elif x <= quartiles[0.5]:
        return 2
    elif x <= quartiles[0.75]:
        return 3
    else:
        return 4

# Assign scores to the RFM values
rfm_data['R'] = rfm_data['Recency'].apply(assign_r_score, args=(quartiles,))
rfm_data['F'] = rfm_data['Frequency'].apply(assign_fm_score, args=(quartiles['Frequency'],))
```

```
rfm_data['M'] = rfm_data['Monetary'].apply(assign_fm_score, args=(quartiles['Monetary'],))
```

## RFM Score Calculation:

Combine the R, F, and M scores to calculate the RFM score for each customer.

```
In [60]: # Calculate the RFM score by combining R, F, and M
rfm_data['RFM_Score'] = rfm_data['R'] * 100 + rfm_data['F'] * 10 + rfm_data['M']
```

```
In [66]: rfm_data['RFM_Score']
```

```
Out[66]: _CustomerID
1      221
2      211
3      144
4      333
5      132
6      412
7      422
8      311
9      244
10     133
11     244
12     444
13     444
14     321
15     411
16     412
17     244
18     244
19     433
20     232
21     434
22     411
23     133
24     111
25     333
26     122
27     411
28     211
29     444
30     433
31     421
32     444
33     424
34     144
35     112
36     422
37     123
38     211
39     444
40     213
41     432
42     433
43     312
44     422
45     323
46     421
47     242
48     143
49     221
50     133
Name: RFM_Score, dtype: int64
```

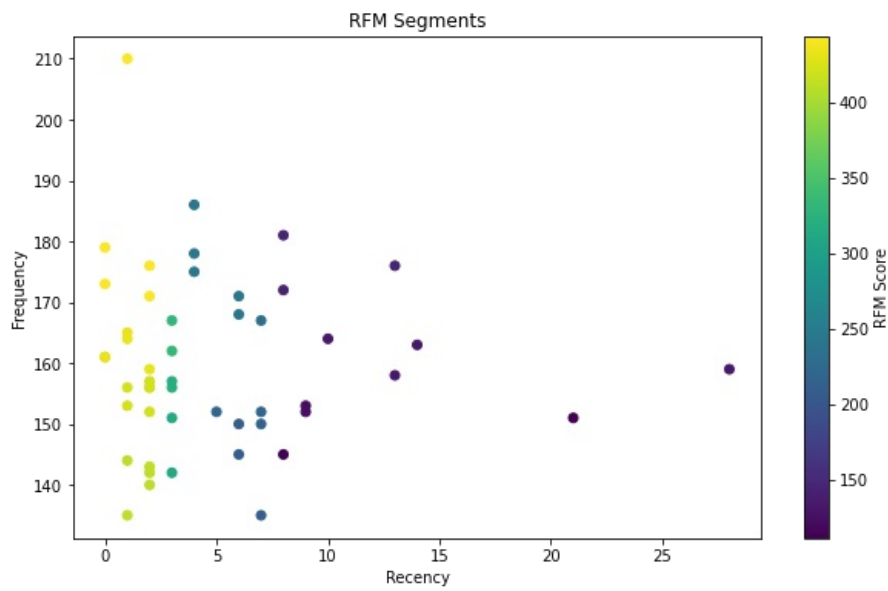
```
In [67]: rfm_data['RFM_Score'].max()
```

```
Out[67]: 444
```

```
In [68]: rfm_data['RFM_Score'].min()
```

```
Out[68]: 111
```

```
In [62]: # Visualize RFM Segments
plt.figure(figsize=(10, 6))
plt.scatter(rfm_data['Recency'], rfm_data['Frequency'], c=rfm_data['RFM_Score'], cmap='viridis')
plt.xlabel('Recency')
plt.ylabel('Frequency')
plt.title('RFM Segments')
plt.colorbar(label='RFM Score')
plt.show()
```



This provides a visual representation of how customers are distributed based on their RFM scores.

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js